PROMPT H26 - Milestone 2: Project Brief

Student: Adhurim Thaqi **Project:** Modality 1 (Text) - Bachelor Thesis Meeting Prep Coach **Date:** 12.10.2026 **Format:** Individual project

---

1 · Stakeholder and target audience

Primary user: an HSLU bachelor student preparing for a meeting with their thesis supervisor - typically holding 20-30 minutes of the supervisor's time and a half-formed list of worries.

Secondary beneficiary: the supervisor, who gets a meeting with a structure and a written follow-up instead of a drifting conversation.

Per the project brief, the models are simulated scenarios: no real supervisor or student is contacted, and the toolkit is an educational prototype, not official HSLU communication, policy or advice.

---

2 · The real-world need

Supervision time is scarce and unevenly used. Students arrive with an unstructured set of concerns, so:

- the meeting drifts and runs out of time
- the most important question gets asked last, or not at all
- nothing is written down, so the next meeting re-opens settled ground
- procedural uncertainty ("is an extension possible?") eats time that could go to the actual work

None of this is a knowledge problem. The student usually knows what is wrong. The problem is shape: an unordered list of worries is not a plan for a twenty-minute conversation.

---

3 · The task

Take the student's rough, unstructured pre-meeting note and produce:

1. a **meeting agenda** fitted to the stated meeting length
2. **prioritized questions**, most important first
3. a **follow-up email** draft recording what was covered

Plus, shipped with the toolkit: a short **user guide** and a **boundaries statement** written for the student, in plain language.

---

4 · Scope

In scope: preparing for, and following up on, one supervisor meeting.

Out of scope, deliberately:

Not doing	Why
Thesis writing or editing	A different task with different failure modes
Research or methods advice	Belongs to the supervisor
Judging the quality of the student's work	Belongs to the supervisor
Answering HSLU procedure from memory	The model cannot know it, and a confident wrong answer is the worst outcome in this project

The last two are not omissions. They are the design.

---

5 · The prompt / workflow challenge

The brief states it directly: turn a messy student note into a focused, useful meeting plan **without inventing HSLU procedures or replacing supervisor judgement.**

Three failure modes follow.

1 · Invented institutional fact. Asked about deadlines, formatting rules, submission procedure or grading, a model produces something plausible and official-sounding. A student could act on it and miss a real deadline. This is the highest-consequence failure in the project.

2 · Substituted judgement. Asked “is my research question good enough?”, the model should help the student *put that question to the supervisor*, not deliver a verdict that displaces it. Notes often apply pressure here - “just tell me yes or no”.

3 · Invention on thin input. Student notes are incomplete. The model’s default is to close gaps with plausible content rather than flag them, because a generic agenda is easier to produce than an admission that there is not enough to work with.

---

6 · Evaluation plan

Method: 10-case / 80% test (brief option a), supported by comparative blind review (option c). The task has objective pass/fail properties - an invented HSLU deadline is unambiguously a failure - which is what makes the 10-case test appropriate here rather than a purely qualitative rubric.

Corpus. Two separate sets, and keeping them separate is what makes the test measure anything:

- **three shipped scenarios** - the product, tuned against
- **ten synthetic test notes** - the evaluation set, never tuned against

Edge cases built into the test set deliberately: a direct procedural question; a demand for a verdict; a three-word note; a note that contradicts itself; a note raising a personal difficulty alongside the academic one.

Hard fails - one occurrence fails the case:

1. States an HSLU procedure, deadline, regulation or grading rule as fact
2. Delivers a verdict that belongs to the supervisor
3. Invents a concern not present in the note
4. Silently drops a concern that was in the note
5. Alters a fact stated in the note

Scored criteria: questions genuinely prioritized rather than listed; agenda fits the stated meeting length; gaps surfaced as open questions; follow-up email accurately reflects the agenda; tone appropriate for writing to a supervisor; usable without heavy editing.

Baseline comparison: one simple out-of-the-box prompt, then at least two meaningful refinements, across the same ten cases.

Three runs per case. W2 established that sampling makes the same prompt produce different answers. A constraint respected twice and broken on the third attempt has not been respected; a single run would measure luck alongside prompt quality.

Independent review: two or three HSLU students who have had supervisor meetings score a sample of outputs blind, without being told which version came from which prompt. This addresses the risk, named in the brief, of evaluating only my own preferred outputs in an individual project.

---

7 · Ethical risks and responsible use

Risk	Response
Invented institutional fact - a student misses a real deadline acting on a confident wrong answer	Never answered from model memory. Either grounded in the real regulations with a citation, or refused and redirected to the official source
Displacing the supervisor - the student takes the tool’s verdict instead of asking	Judgement questions are structurally rerouted into questions the student asks. Stated to the user in the boundaries statement, not only in the report
Mistaken for official HSLU material	Labelled an educational prototype in the output itself
Privacy - real notes describe identifiable supervisors and personal circumstances	Synthetic test notes only. No real note, supervisor or meeting is used
Wellbeing - a note may raise a personal difficulty	Not coached on. Routed to HSLU student support services
Accessibility	Plain-text, screen-reader-friendly output; user guide assumes no prior familiarity with AI tools

Disclosure: full AI-use disclosure in the individual process record.

---

8 · One open design decision

W4 introduced retrieval, which changes the options for failure mode 1.

Milestone 1 committed to **refusing** every institutional question, on the grounds that an ungrounded model asked about a deadline can only guess. With RAG there is a third option: **ground** the toolkit in the real HSLU regulations so it answers what is documented, cites it, and refuses only what lies outside the documents.

Grounding is the better product if a usable document set exists. It is also the more interesting evaluation, because it splits one refusal into two:

- *not in my documents* - go and check the official source
- *no document could answer this* - it is your supervisor's judgement

Those are different answers and a toolkit that conflates them is worse than one that separates them.

To be settled by Milestone 3, once it is clear what documentation is actually available. The refusal behaviour is kept either way; the question is only how wide its scope is.

---

9 · Deliverables

Deliverable	Form
Three realistic student scenarios	Written cases, with agenda + questions + follow-up email for each
Prototype	Python chain following the W4 pattern, runnable in Colab
User guide	Short, plain language
Boundaries statement	What the toolkit will not do, written for the student
Evidence	Prompt log (verbatim prompts, every version), run log, scored test results
Report + process record	Per the module deliverables sheet

---

10 · Feasibility

The W3 criteria for a feasible idea were a defined user, a specific task, and a product that can realistically be built.

- **Defined user:** a bachelor student with a supervisor meeting this week.
- **Specific task:** one note in, three artifacts out.
- **Realistic product:** text only, no interface required beyond a notebook, and the evaluation runs on ten synthetic cases I write myself. Nothing in the plan depends on recruiting participants or on access I do not have.

The main risk is not technical. It is leaving the test runs until late, when the iteration evidence needs weeks of accumulation to be worth anything.