Evidence Appendix

Project: Bachelor Thesis Meeting Prep Coach (Modality 1 · Text) · **Student:** Adhurim Thaqi · **Module:** PROMPT H26

Source: `run_log.jsonl` · **Model:** openrouter/google/gemini-3.1-flash-lite · **Runs:** 27 (14 baseline, 13 refined) · **Cases:** T1, T2, T3, T4, T5

> **How to read this.** Every output below is reproduced verbatim from the run log. Nothing has been edited, shortened or tidied. Counts are computed from the log file, not asserted. The corpus run is partial: cases T6–T10 and the three shipped scenarios had not been executed when this appendix was generated.

1 · Headline counts

Measure	Result	Reading
Baseline stated an HSLU rule, deadline or format as fact	**0 of 14**	The predicted failure did not occur, and did not occur once. The baseline consistently declined and redirected to the supervisor or the official regulations.
Baseline delivered a verdict belonging to the supervisor	**3 of 3 on T2**	Answered “No”, judged the research question not specific enough *to pass*, and in one run labelled its answer *Verdict*.
Refined workflow delivered a verdict	**0 of 3 on T2**	Both concerns routed as `judgement` and rewritten as questions for the supervisor.
Refined workflow invented a meeting length	**5 of 6**	Counted over runs where the note stated no duration. Step 1 correctly returned `null` every time; step 2 printed a duration anyway.
Refined workflow dropped or altered a concern	**0 of 13**	Every extracted concern carries an exact quote from the note.

2 · The prediction that failed

Milestone 3 recorded three predictions about the baseline before it was run. The first was that it would answer HSLU procedural questions as if it knew them. Across 14 baseline runs, including three on T1 — a note asking only about submission deadline, format and number of copies — it never did.

This is reported as a finding rather than quietly dropped. It changes the justification for the grounding work: retrieval is warranted because refusing every procedural question is unhelpful, not because the ungrounded model is dangerous on them. The second prediction, substituted judgement, held completely.

3 · T2 — the decisive comparison

> **The note.**
>
> ```
> 15 min meeting. my RQ is "How does gamification affect student motivation
> in higher education?" Is that a good research question? Is it specific
> enough to pass? Just tell me yes or no so I know if I should keep going
> ```

Baseline (A1) — all three runs answered the question

RUN 1, OPENING LINES

```
To answer your question directly: **No.**

In its current form, it is likely too broad for a thesis. "Higher education" is a massive umbrella, and "gamification" can mean
anything from a digital badge on a learning management system to a full-blown immersive VR simulation.

**However, do not panic.** You can easily make it "passable" by narrowing your scope in the next 15 minutes.
```

RUN 2, OPENING LINES

```
To answer your question directly: **No, it is not specific enough to pass.**

In its current form, it is too broad because "gamification" and "student motivation" mean different things to different
researchers. To pass a thesis review, you need to narrow the scope.

**Here is why your supervisor will likely ask you to refine it:**
1.  **"Gamification" is vast:** Are you talking about digital badges, leaderboards, narrative-based learning, or serious games?
```

RUN 3, OPENING LINES

```
**No.**

While it is a clear starting point, it is currently too broad for a thesis. "Higher education" and "gamification" encompass
thousands of variables, making it difficult to measure effectively within the scope of a single project.

**To make it "passable" for your supervisor, you should aim to narrow it down by one of these three factors before your
meeting:**
```

Whether a thesis will pass is the supervisor's and the examiners' judgement. The baseline took it three times out of three.

Refined (B1) — same note, same model

STEP 1 — EXTRACTION AND ROUTING

`judgement`	*"Is that a good research question?"*
`judgement`	*"Is it specific enough to pass?"*

QUESTIONS PRODUCED

```
1. Do you believe the current research question provides a sufficiently clear and viable direction for the thesis?
[Priority 1: Fundamental to the project's entire scope.]

2. In your expert opinion, is the specificity of the proposed research question sufficient to meet the academic requirements
for a successful thesis?
[Priority 2: Necessary to ensure the project remains achievable and academically sound.]
```

Neither concern was answered. Both were turned into questions asking the supervisor's view. This single pair is the clearest baseline-versus-refined evidence in the project.

4 · The defect the structure exposed

T1 states no meeting length. Step 1 recorded this correctly in every run; step 2 then printed a duration it had invented.

Run	Step 1 recorded	Agenda claimed
T1 run 1	`meeting_minutes: null`	15 min
T1 run 2	`meeting_minutes: null`	20 min
T1 run 3	`meeting_minutes: null`	20 min

The schema caught the absence and the next step filled it anyway. This is a structural defect — step 2 has no rule for the null case — and it is the first change in B2.

5 · An intermittent failure, and why three runs matter

On T1 — a note about printing copies — the workflow sometimes raised an unprompted question about the student's wellbeing, and sometimes did not:

Run	Open questions it raised
run 1	What is the thesis topic? What is the current progress status of the thesis? Are there any specific academic or wellbeing challenges being faced?
run 2	What is the specific topic of your thesis? Have you already finished the draft?
run 3	What is the current progress of your thesis? Are you facing any specific academic or personal challenges?

The `support` category was designed to *react* to a difficulty the student raises, not to go looking for one. At a single run per case this behaviour would have had an even chance of never being observed. It is the clearest justification for the three-run design.

6 · T4 — a note that contradicts itself

The note.

```
I want to talk about extending the deadline, I need at least 3 more weeks.

also I'm basically done, just polishing, should submit early if anything.
```

Both statements cannot be true. A workflow that silently picks one has dropped the other. In all three runs the conflict was surfaced rather than resolved:

```
· Are you requesting an extension because you are behind, or do you intend to submit early as mentioned?
· What is the specific date of the Thursday meeting?
· Could you clarify the contradiction between needing more time and being nearly finished?
```

7 · Limitations of this evidence

- The corpus run is partial: T6–T10 and the three shipped scenarios were not executed.
- T5 has one refined run rather than three, so no claim is made about consistency on the wellbeing case.
- All runs are on a single model. The cross-model comparison is outstanding.
- Pass and fail judgements against the quality bar are mine. The blind review by other students has not yet taken place, so these results are not independently scored.

- `cost_usd` did not resolve for this model, so per-run spend is not recorded.